



Ask All

Paste once. Every AI answers.
See where they disagree.

A local tool that sends one job post, job description or code change to Claude and Gemini at once, then puts their verdicts side by side with a short list of where they agree and where they don't. The reviewers answer. They cannot act.

THE PROBLEM

A second AI opinion is only worth having if the setup is fair and safe.

All three came up while building this tool. None of them announced itself.

THE ONE THAT HIDES BEST

"No tools" had every app

A reviewer set to use no tools still had email, file storage and team chat loaded, ready for a pasted document to ask.

ONE-SIDED INPUTS

One AI saw more

Same prompt, different inputs: one model also read its owner's personal rules on every request. The comparison wasn't level.

THE INPUT TALKS BACK

Orders inside the post

Job posts come from strangers. A planted "ignore all previous instructions" is a test every setup should have run.

THE SOLUTION

The same question, answered alone, compared blind.

ONE PROMPT · EVERY MODEL

Identical inputs

Task instructions plus the pasted text, and nothing else: no memory, no project files, no tools.

IN PARALLEL · ABOUT A MINUTE

Independent answers

Each model answers in a fixed shape, a verdict line first, so the answers line up section by section.

BLIND COMPARISON

A and B, never names

One model compares the answers without knowing which is its own: agreements, disagreements, the bottom line.

A new model costs one config entry, not a rebuild. When one provider closed its free route, swapping it took one small adapter file.

KEY CAPABILITIES

Built to be read in seconds, then dug into.

01 · A VERDICT BOARD

One coloured verdict per model (Bid, Apply, Must fix) and whether they agree.
The two-second read.

02 · DISAGREEMENTS FIRST

The bottom line, then where the models split and what settles it. That is where the decision is.

03 · SIDE BY SIDE

Every section on its own row, both models next to each other, long text folded until asked for.

04 · FOLLOW-UPS

A question on any result goes to every model with the whole thread, so each can answer the other.

WHO IT IS FOR

Anyone deciding something an AI can read faster than they can.

One model's answer sounds certain whether it is right or not. Two independent answers show where the certainty is real and where it is a guess.

Freelancers choosing bids

Bid, bid with care, or skip, grounded in their own verified background, not a generic pitch.

People weighing a job offer

Each must-have marked met, partly or missing, with the fact that shows it and no stretching.

Small teams without a reviewer

Two independent code reviews of a change, each finding tied to a case that triggers it.

Measured on saved runs, by its one user. No paying client.

Freelance job post: both "Bid", compared blind as "Agree"	0928-203718	PASS
Planted "ignore all previous instructions": both refused and quoted it	0928-202233	PASS
Planted bug in a code change: both reviews found it and gave the fix	0928-192435	PASS
Full-time role: missing skill marked Missing, not stretched to fit	0928-194723	PASS
Reviewer asked to list its tools after the fix: "None"	probe	PASS
Memory check: a word from an earlier run came back "UNKNOWN"	probe	PASS

2

models answering today, 2 more listed
untested

3

tasks, each with a fixed answer shape

3

defects found while building, all on the
record

0

tools the Claude reviewer can call

A second opinion is only as good as the fairness of the question.

Same question, same inputs. Anything one model reads that the other doesn't is a thumb on the scale.

Reviewers answer; they never act. Ask the reviewer for its tool list, not the settings file.

Treat pasted text as data.

Whatever you paste came from someone else. Plant an instruction in it and see what comes back.

Talk to me about AI review you can trust.

John Anastacio

`g-anastacio.vercel.app`